

Below is the take-home assignment we'd like you to complete as part of our hiring process. Please read it carefully and submit your solution according to the guidelines provided.

## Take-Home Assignment: LLM-Based Ticket Reply Evaluation

### Objective

You are given a CSV file (tickets.csv) containing two columns:

1. ticket – A (fictional) customer support ticket message.
2. reply – A response generated by an AI system.

Your task is to **use an OpenAI Large Language Model (LLM)** (e.g., GPT-4o or o1) to **evaluate** each reply along two dimensions:

- **Content** (relevance, correctness, completeness)
- **Format** (clarity, structure, grammar/spelling)

You will then produce a **new CSV** containing four additional columns:

1. content\_score
2. content\_explanation
3. format\_score
4. format\_explanation

Please use the scoring scale **1-5** for the “\_score” fields and provide a **short textual explanation** in each of the “\_explanation” fields.

### Assignment Requirements

#### Input

A CSV file named tickets.csv with the columns:

- ticket (the customer's message)
- reply (the AI-generated response)

#### Processing with an LLM

You will write a Python script or Jupyter notebook that:

1. Reads tickets.csv.

2. For each row, sends both ticket and reply to an LLM with a suitable prompt asking the model to evaluate the reply's content and format.
3. Parses the LLM's output and extracts two numeric scores (content\_score, format\_score) and two explanations (content\_explanation, format\_explanation).
4. Feel free to engineer your prompt in a way that you find most effective.
5. Ensure that your code can handle potential errors (e.g., missing data or API issues).

## Output

A CSV file named tickets\_evaluated.csv with the following columns:

- ticket
- reply
- content\_score
- content\_explanation
- format\_score
- format\_explanation

## Deliverables

### Notebook or Script

Provide a Jupyter notebook (.ipynb) or a Python script (.py) containing your solution.

### tickets\_evaluated.csv

The final output with the added scoring columns.

### README

- Briefly describe how to set up and run your code.
- Mention any dependencies, environment variables (e.g., for API keys), or library installations needed.

### (Optional but Encouraged) Tests

If you have time, include a small test suite (e.g., using pytest or unittest) that covers your core functions (reading CSV, calling the LLM, writing output, etc.).

### Handling API Keys

Please do **not** commit any private keys to your repository or submission.

Use environment variables or a .env file (excluded from version control) for storing sensitive credentials.

## Development Guidelines

Write **clean, maintainable, and well-documented** Python code.

You can use any packages you prefer; just ensure they are listed (e.g., in requirements.txt or pyproject.toml).

Explanations can be brief (one or two sentences).

## Submission

Send us a link to a public GitHub repository, a zipped folder, or any other method that allows us to easily review your code.

Include instructions in your README on how to run the solution.

## Evaluation Criteria

### Code Quality

Readability, structure, and adherence to best practices (e.g., PEP-8, error handling).

### Prompt Engineering

How effectively you prompt the LLM to produce useful scores and explanations.

### Correctness of Output

Your approach to reading/writing CSVs, generating valid numeric scores, and capturing explanations.

### Documentation & Tests

Clarity of your README and any test coverage you provide.

We expect this assignment to take about **2–4 hours** in total, depending on your familiarity with LLM APIs. If you have any questions or need clarifications, feel free to reach out.

**Good luck, and we look forward to seeing your solution!**